

Deep learning in R using MXNet

Qiang Kou

qkou@umail.iu.edu

R/Finance 2016, Chicago



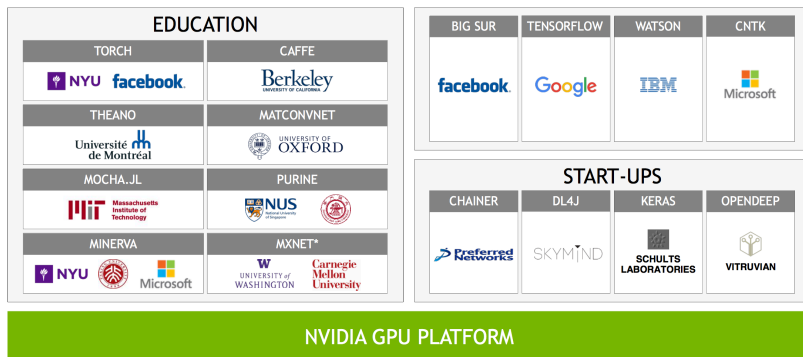
INDIANA UNIVERSITY

- Qiang Kou, KK, 寇强
- PhD student in bioinformatics from Indiana University
- Rcpp team member

What is MXNet?

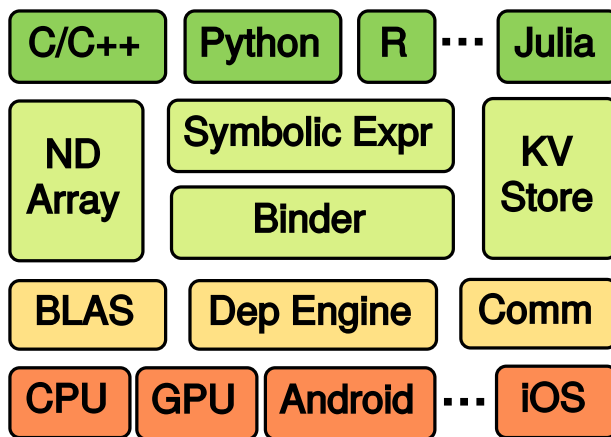
Deep learning platforms

THE ENGINE OF MODERN AI



<http://www.slideshare.net/NVIDIA/nvidia-ces-2016-press-conference>

MXNet



MXNet: A Flexible and Efficient Machine Learning Library for Heterogeneous Distributed Systems

Why MXNet?

Windows support

```

RGui (64-bit) - [R Console]
File Edit View Misc Packages Windows Help

> drat::addRepo("dmlc")
> install.packages("mxnet")
--- Please select a CRAN mirror for use in this session ---
trying URL 'http://dmlc.github.io/drat/bin/windows/contrib/3.2/mxnet_0.5.zip'
Content type 'application/zip' length 13372314 bytes (12.8 MB)
downloaded 12.8 MB

package 'mxnet' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  C:\Users\Administrator\AppData\Local\Temp\2\Btmpc2G04H\downloaded_packages
> require("mxnet")
Loading required package: mxnet
> require("mlbench")
Loading required package: mlbench
> data(Sonar, package="mlbench")
> Sonar[,61] = as.numeric(Sonar[,61])-1
> train.ind = c(1:50, 100:150)
> train.x = data.matrix(Sonar[train.ind, 1:60])
> train.y = Sonar[train.ind, 61]
> test.x = data.matrix(Sonar[-train.ind, 1:60])
> test.y = Sonar[-train.ind, 61]
> mx.set.seed(0)
> model <- mx.mlp(train.x, train.y, hidden_node=10, out_node=2, out_activation="softmax",
+               num.round=20, array.batch.size=15, learning.rate=0.07, momentum=0.9,
+               eval.metric=mx.metric.accuracy)
Auto detect layout of input matrix, use rowmajor..
Start training with 1 devices
[1] Train-accuracy=0.488888888888889
[2] Train-accuracy=0.514285714285714
[3] Train-accuracy=0.514285714285714
[4] Train-accuracy=0.514285714285714

```

Installation

DRAT repo

For Windows and Mac users

```
install.packages("drat", repos="https://cran.rstudio.com")  
drat::addRepo("dmlc")  
install.packages("mxnet")
```

MNIST demo

MNIST demo

```
library(mxnet)
train <- read.csv("train.csv", header=TRUE)
test <- read.csv("test.csv", header=TRUE)
train <- data.matrix(train)
test <- data.matrix(test)
train.x <- train[,-1]
train.y <- train[,1]
train.x <- t(train.x/255)
test <- t(test/255)
```

The dataset: <https://www.kaggle.com/c/digit-recognizer/data>

MNIST demo

```
data <- mx.symbol.Variable("data")
fc1 <- mx.symbol.FullyConnected(data, name="fc1", num_hidden=128)
act1 <- mx.symbol.Activation(fc1, name="relu1", act_type="relu")
fc2 <- mx.symbol.FullyConnected(act1, name="fc2", num_hidden=64)
act2 <- mx.symbol.Activation(fc2, name="relu2", act_type="relu")
fc3 <- mx.symbol.FullyConnected(act2, name="fc3", num_hidden=10)
softmax <- mx.symbol.SoftmaxOutput(fc3, name="sm")
```

MNIST demo

```
graph.viz(softmax$as.json())
```

Computation graph



MNIST demo

```
mx.set.seed(0)
model <- mx.model.FeedForward.create(softmax, X=train.x,
                                     y=train.y, ctx=mx.gpu(), num.round=10,
                                     array.batch.size=100,
                                     learning.rate=0.07, momentum=0.9,
                                     eval.metric=mx.metric.accuracy,
                                     initializer=mx.init.uniform(0.07),
                                     batch.end.callback
                                     = mx.callback.log.train.metric(100))
```

MNIST demo

```

tmux
Batch [300] Train-accuracy=0.975033333333334
Batch [400] Train-accuracy=0.976425000000003
[4] Train-accuracy=0.976595238095241
Batch [100] Train-accuracy=0.9795
Batch [200] Train-accuracy=0.979499999999999
Batch [300] Train-accuracy=0.980900000000001
Batch [400] Train-accuracy=0.982150000000003
[5] Train-accuracy=0.982500000000003
Batch [100] Train-accuracy=0.9867
Batch [200] Train-accuracy=0.9856
Batch [300] Train-accuracy=0.985900000000001
Batch [400] Train-accuracy=0.986325000000003
[6] Train-accuracy=0.986404761904765

| 0 GeForce GTX 960 Off | 0000:01:00.0 On | N/A |
| 22% 38C P2 48W / 160W | 641MiB / 2044MiB | 50% Default |
+-----+-----+-----+
+-----+-----+-----+
| Processes: | GPU Memory |
| GPU PID Type Process name Usage |
+-----+-----+-----+
| 0 1214 G /usr/bin/X 181MiB |
| 0 1917 G compiz 98MiB |
| 0 2828 G /usr/lib/firefox/firefox 1MiB |
| 0 25748 C /usr/lib/R/bin/exec/R 346MiB |
+-----+-----+-----+
+ Downloads
[3] 0:~/Downloads* "gkou-All-Series" 21:12 22-Mar-16

```

Convolutional Neural Network

```
# input
data <- mx.symbol.Variable('data')
# first conv
conv1 <- mx.symbol.Convolution(data=data, kernel=c(5,5), num_filter=20)
tanh1 <- mx.symbol.Activation(data=conv1, act_type="tanh")
pool1 <- mx.symbol.Pooling(data=tanh1, pool_type="max", kernel=c(2,2), stride=c(2,2))
# second conv
conv2 <- mx.symbol.Convolution(data=pool1, kernel=c(5,5), num_filter=50)
tanh2 <- mx.symbol.Activation(data=conv2, act_type="tanh")
pool2 <- mx.symbol.Pooling(data=tanh2, pool_type="max", kernel=c(2,2), stride=c(2,2))
# first fullc
flatten <- mx.symbol.Flatten(data=pool2)
fc1 <- mx.symbol.FullyConnected(data=flatten, num_hidden=500)
tanh3 <- mx.symbol.Activation(data=fc1, act_type="tanh")
# second fullc
fc2 <- mx.symbol.FullyConnected(data=tanh3, num_hidden=10)
# loss
lenet <- mx.symbol.SoftmaxOutput(data=fc2)
```


LSTM

```
# lstm cell symbol
lstm <- function(num.hidden, indata, prev.state, param, seqidx, layeridx, dropout=0) {
  if (dropout > 0)
    indata <- mx.symbol.Dropout(data=indata, p=dropout)
  i2h <- mx.symbol.FullyConnected(data=indata,
                                   weight=param$i2h.weight,
                                   bias=param$i2h.bias,
                                   num.hidden=num.hidden * 4,
                                   name=paste0("t", seqidx, ".l", layeridx, ".i2h"))
  h2h <- mx.symbol.FullyConnected(data=prev.state$h,
                                   weight=param$h2h.weight,
                                   bias=param$h2h.bias,
                                   num.hidden=num.hidden * 4,
                                   name=paste0("t", seqidx, ".l", layeridx, ".h2h"))


  gates <- i2h + h2h
  slice.gates <- mx.symbol.SliceChannel(gates, num.outputs=4,
                                       name=paste0("t", seqidx, ".l", layeridx, ".slice"))

  in.gate <- mx.symbol.Activation(slice.gates[[1]], act.type="sigmoid")
  in.transform <- mx.symbol.Activation(slice.gates[[2]], act.type="tanh")
  forget.gate <- mx.symbol.Activation(slice.gates[[3]], act.type="sigmoid")
  out.gate <- mx.symbol.Activation(slice.gates[[4]], act.type="sigmoid")
  next.c <- (forget.gate * prev.state$c) + (in.gate * in.transform)
  next.h <- out.gate * mx.symbol.Activation(next.c, act.type="tanh")

  return (list(c=next.c, h=next.h))
}
```

A shiny app


A shiny app



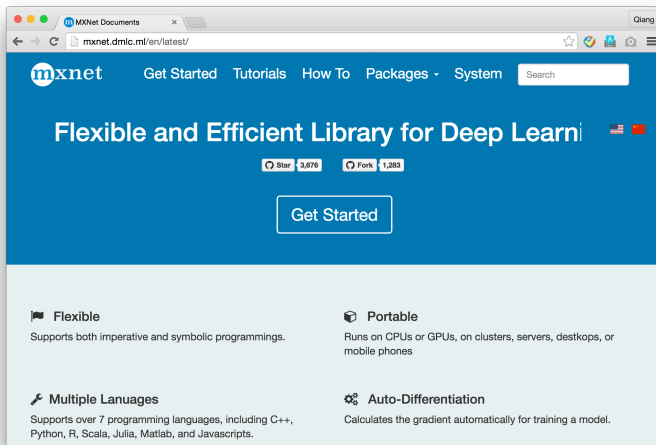
A shiny app

```
model <- mx.model.load("Inception/Inception_BN", iteration = 39)
synsets <- readLines("Inception/synset.txt")
mean.img <-
  as.array(mx.nd.load("Inception/mean_224.nd")[["mean_img"]])
im <- load.image(src)
normed <- preproc.image(im, mean.img)
prob <- predict(model, X = normed)
```

Acknowledgment



Go to <http://mxnet.dmlc.ml/> to get started!



Thank you for the time!